

# Lab2

## Pthread & OpenMP

18 Oct 2018  
Parallel Programming

# SLURM quick reference

---

```
srun [flags] ./prog
```

```
===== or =====
```

```
#!/bin/bash
```


```
#SBATCH [flags]
```

```
srun ./prog # (MPI)
```

```
./prog # (non-MPI)
```

```
----- run with: -----
```

```
sbatch job.sh
```



[flags]:

- p debug or batch
- N number of nodes
- n number of processes
- c **CPUs per process**
- t additional time limit
- J name of job

# Outline

— — —

- Pthread
  - Hello world
  - Mutex
  - Condition variable
- OpenMP
- OpenMP + MPI

# Running pthread programs on apollo

— — —

## SYNOPSIS

```
#include <pthread.h>
```

```
int pthread_create(  
    pthread_t *thread, const pthread_attr_t *attr,  
    void *(*start_routine) (void *), void *arg);
```

Compile and link with -pthread.

Type `man pthread_create` in terminal to see this

# Running pthread programs on apollo

— — —

cp /home/ta/lab2/hello\_pthread.c .

compile

gcc hello\_pthread.c -o hello\_pthread -pthread

execute

srun -c4 -n1 ./hello\_pthread 4

-c4 means 4 CPUs per process

-n1 means 1 process

You can use sbatch as well



NOT  
-lpthread

# Outline

— — —

- **Pthread**
  - Hello world
  - **Mutex**
  - Condition variable
- OpenMP
- OpenMP + MPI

# Pthread Lock/Mutex Routines

- To use mutex, it must be declared as of **type pthread\_mutex\_t** and initialized with **pthread\_mutex\_init()**
- A mutex is destroyed with **pthread\_mutex\_destroy()**
- A critical section can then be protected using **pthread\_mutex\_lock()** and **pthread\_mutex\_unlock()**
- Example:

```
#include "pthread.h"
pthread_mutex_t mutex;
pthread_mutex_init (&mutex, NULL);
pthread_mutex_lock(&mutex);

    Critical Section

pthread_mutex_unlock(&mutex);
pthread_mutex_destroy(&mutex);
```

specify default attribute for the mutex

// enter critical section

// leave critical section

# Mutex

— — —

man pthread\_mutex\_init

```
#include <pthread.h>
```

```
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
```

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

man pthread\_mutex\_lock



# Mutex:

## [Practice 1] approximate $\pi$ using pthread

---

```
gcc pi_pthread.c -o pi_pthread -pthread -lm
```

code  
filename

executable  
filename

we're using  
pthread

```
srunk -c4 -n1 ./pi_pthread 4 500000000
```

number of  
threads

number of  
points



# Outline

— — —

- **Pthread**
  - Hello world
  - Mutex
  - **Condition variable**
- OpenMP
- OpenMP + MPI

# Condition Variables (CV)

- CV represent some **condition** that a thread can:

— — —

- Wait on, until the condition occurs; or
- Notify other waiting threads that the condition has occurred

- Three operations on condition variables:

- **wait()** --- **Block** until another thread calls **signal()** or **broadcast()** on the CV
- **signal()** --- Wake up **one thread** waiting on the CV
- **broadcast()** --- Wake up **all threads** waiting on the CV

- In Pthread, CV **type** is a **pthread\_cond\_t**

- Use **pthread\_cond\_init()** to initialize
- **pthread\_cond\_wait (&theCV, &somelock)**
- **pthread\_cond\_signal (&theCV)**
- **pthread\_cond\_broadcast (&theCV)**

# Condition variable

— — —

```
#include <pthread.h>
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;
int pthread_cond_signal(pthread_cond_t *cond);
int pthread_cond_broadcast(pthread_cond_t *cond);
int pthread_cond_wait(
    pthread_cond_t *restrict cond,
    pthread_mutex_t *restrict mutex);
```

man pthread\_cond\_broadcast

man pthread\_cond\_timedwait

# Condition variable

— — —

## `pthread_cond_signal`

Make 1 thread that called `pthread_cond_wait` to continue

## `pthread_cond_broadcast`

Make all threads that called `pthread_cond_wait` to continue

## `pthread_cond_wait`

Wait for some other thread to call `pthread_cond_signal`

# [Practice 2] Condition variable

---

Compile and execute `pthread_cond.c`  
under `lab2/`

you should see →  
which is incorrect

Modify the program using condition  
variable so it will wait until you input  
the values then print



Threads have been created

Enter 4 values

Values filled in array are

0

0

0

0

# Outline

— — —

- Pthread
  - Hello world
  - Mutex
  - Condition variable
- **OpenMP**
- OpenMP + MPI

# Running OpenMP programs on apollo: example (/home/ta/lab2)

---

compile

```
gcc hello_omp.c -o hello_omp -fopenmp
```

execute


```
srun -c4 -n1 ./hello_omp
```

-c4 means 4 CPUs per process

-n1 means 1 process

You can use sbatch as well

Try different number of threads!



OpenMP automatically  
detects number of CPUs  
from SLURM (affinity)  
So we don't have to  
specify it again



# Count prime numbers: sequential version

---

```
gcc -lm prime.c -o prime
```

```
srun ./prime 1000
```

```
srun ./prime 10000000
```

# Count prime numbers:

## [Practice 3] OpenMP


---

1. Modify the sequential prime.c with openmp
2. Try to see the effect of changing  
dynamic/static **scheduling**  
**chunk size**  
number of **threads**

[example commands]

```
gcc -lm prime_omp.c -o prime_omp -fopenmp
```

```
srun -c4 -n1 ./prime_omp 10000000
```



# Outline

— — —

- Pthread
  - Hello world
  - Mutex
  - Condition variable
- OpenMP
- **OpenMP + MPI**

# Hybrid MPI and OpenMP program

---

```
mpicc hello_hybrid.c -o hello_hybrid -fopenmp
```

We're  
using MPI

We're using  
OpenMP

```
srun -c3 -n2 ./hello_hybrid
```

3 threads

2  
processes

# Hybrid MPI and OpenMP program: Hello World

---

```
srun -c 3 -n2 -N2 ./hello_hybrid
```

```
Hello apollo32: rank  0/ 2, thread  0/ 3
```

```
Hello apollo32: rank  0/ 2, thread  1/ 3
```

```
Hello apollo32: rank  0/ 2, thread  2/ 3
```

```
Hello apollo33: rank  1/ 2, thread  0/ 3
```

```
Hello apollo33: rank  1/ 2, thread  1/ 3
```

```
Hello apollo33: rank  1/ 2, thread  1/ 3
```

# Hybrid MPI and OpenMP program:

## **[Practice 4]** Approximate $\pi$

---

Use MPI and OpenMP to approximate  $\pi$

(You can refer to your lab1b code)

