

# Parallel Programming

---

Lab4 2018/12/6

# Overall Performance Optimization Strategies

- ❖ Maximize parallel execution to achieve maximum utilization
- ❖ Optimize memory usage to achieve maximum memory throughput
- ❖ Optimize instruction usage to achieve maximum instruction throughput

# CUDA Optimization technique

- ❖ Resource information & device information
- ❖ 2D data placement
- ❖ Share memory
- ❖ Streaming

# Resource information & device information

❖ Device information -> deviceQuery

❖ Resource information

- add “-Xptxas=-v” to compiler flag to check the resource allocation
- add “-arch=sm\_61” to make the compiler allocate suitable resource (default is sm\_30)

```
[zls80826@hades01 lab4]$ make
nvcc -O2 -std=c++11 -Xptxas=-v -arch=sm_61 -lpng -lz -o sobel sobel.cu
ptxas info      : 50 bytes gmem
ptxas info      : Compiling entry function '_Z5sobelPhS_mmjjj' for 'sm_61'
ptxas info      : Function properties for _Z5sobelPhS_mmjjj
    0 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
ptxas info      : Used 32 registers, 590 bytes smem, 364 bytes cmem[0], 24 bytes cmem[2]
```


# Occupancy

- ❖ With the above information, we can calculate how many blocks can concurrently run on a SM
- ❖ Number of active blocks are limited by
  - Shared memory
  - Register
  - Max threads/threads per block
- ❖ We can calculate the occupancy manually or API
  - `__host__ __device__ cudaError_t cudaOccupancyMaxActiveBlocksPerMultiprocessor ( int* numBlocks, const void* func, int blockSize, size_t dynamicSMemSize )`
  - [Reference](#)

# 2D Data placement


- ❖ Most of time we use GPU devices to solve the 2D, 3D problems
- ❖ Official document suggest using 2D API for 2D memory operation
- ❖ Use `cudaMallocPitch` to allocate memory
- ❖ `cudaMemcpy2D` for memory copy

```
T* pElement = (T*)((char*)BaseAddress + Row * pitch) + Column
```



# Share memory

- ❖ Share memory is faster than device memory
- ❖ Share memory take advantages on the applications with data locality
- ❖ Example for sobel.



Be accesed for 3 times

# Streaming

## ❖ Reference

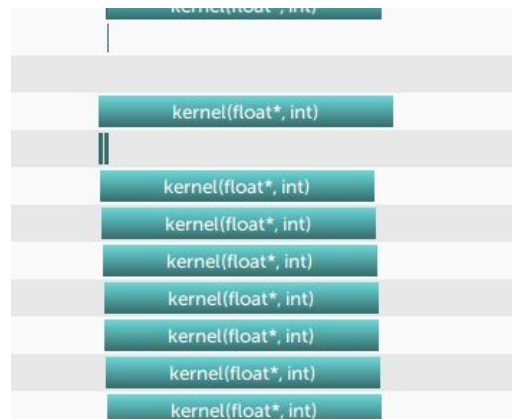
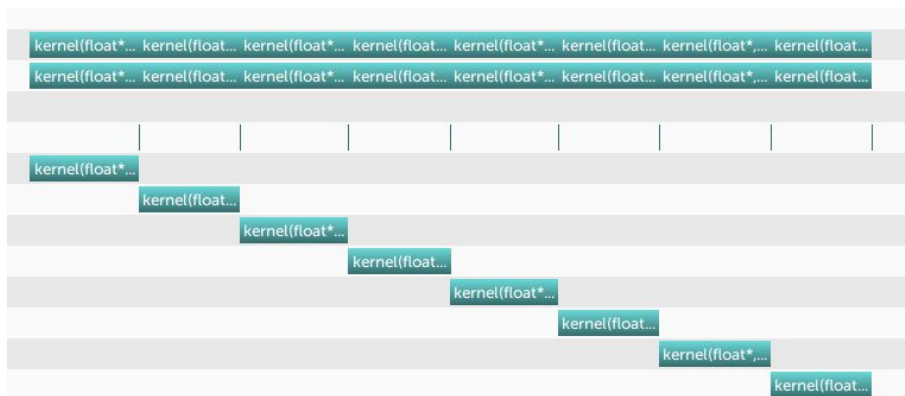
```
--default-stream {legacy|null|per-thread} (-default-stream)
Specify the stream that CUDA commands from the compiled program will be sent
to by default.

legacy
    The CUDA legacy stream (per context, implicitly synchronizes with
    other streams).

per-thread
    A normal CUDA stream (per thread, does not implicitly
    synchronize with other streams).

'null' is a deprecated alias for 'legacy'.

Allowed values for this option: 'legacy','null','per-thread'.
Default value: 'legacy'.
```





# Hints

## ❖ Easy

- Calculate the occupancy and tune the [thread per block]
- Compare the speed
  - mask on share memory, const memory, device memory
- Images are 2D (2D API)
- Take out branch (padding with 0)

## ❖ Hard

- Streaming
- Slides page 7.

0	1	2	3
4	5	6	7
8	9	10	11

		0	1	2	3		
		4	5	6	7		
		8	9	10	11		

# Lab4 Assignment

- ❖ Optimize sobel.cu with any optimization techniques you like
- ❖ We still provide modified cpu code
  - `cp /home/pp18/shared/lab4 ~/homework`
  - `make`
  - `srun -n 1 --gres=gpu:1 -ppp ./sobel <input> <output>`
  - lab4-judge for judging
- ❖ Submission
  - After 12/6 23:59, We only use the scoreboard for scoring
- ❖ Deadline
  - 12/9 23:59

# Grading

	Before 12/6 23:59	Before deadline
Faster than TA	150	125
Faster than <code>benchmark100</code>	100	100
Faster than <code>benchmark60</code>	60	60

Port the sobel on GPU and tune the threadPerBlock

- ❖ In lab(12/6 23:59) only
  - If you can't faster than mewtwo, demo the easy part of the hints can also get 100

# Appendix: nvprof

## metrics

```
[zls80826@hades01 lab4]$ srtn -n 1 --gres=gpu:1 -ppp nvprof --metrics achieved_occupancy ./sobel testcase/large_candy.png out.png
Height: 10800
Width: 16320
Channel: 3
input spend 3.05731 second.
==22404== NVPROF is profiling process 22404, command: ./sobel testcase/large_candy.png out.png
kernel spend 0.504053 second.
output spend 5.66372 second.
==22404== Profiling application: ./sobel testcase/large_candy.png out.png
==22404== Profiling result:
==22404== Metric result:
Invocations      Metric Name      Metric Description      Min      Max      Avg
Device "GeForce GTX 1080 (0)"
  Kernel: sobel(unsigned char*, unsigned char*, unsigned long, unsigned long, unsigned int, unsigned int, unsigned int)
    1      achieved_occupancy      Achieved Occupancy      0.248712      0.248712      0.248712
[zls80826@hades01 lab4]$ nvprof --help^C arep metrics
[zls80826@hades01 lab4]$ srtn -n 1 --gres=gpu:1 -ppp nvprof --metrics sm_efficiency ./sobel testcase/large_candy.png out.png
Height: 10800
Width: 16320
Channel: 3
input spend 3.04934 second.
==22476== NVPROF is profiling process 22476, command: ./sobel testcase/large_candy.png out.png
kernel spend 0.558252 second.
output spend 5.83806 second.
==22476== Profiling application: ./sobel testcase/large_candy.png out.png
==22476== Profiling result:
==22476== Metric result:
Invocations      Metric Name      Metric Description      Min      Max      Avg
Device "GeForce GTX 1080 (0)"
  Kernel: sobel(unsigned char*, unsigned char*, unsigned long, unsigned long, unsigned int, unsigned int, unsigned int)
    1      sm_efficiency      Multiprocessor Activity      98.12%      98.12%      98.12%
```